

# AI Food Recommendation Chatbot: Multi-Agent RAG System

## System Overview

The AI Food Recommendation System is a sophisticated multi-agent conversational platform that combines natural language processing, machine learning, and advanced retrieval techniques to provide personalized food recommendations. The system employs a modular architecture with distributed components working in coordination to deliver contextually relevant suggestions through natural conversation.

## Architecture Components

### *Core System Architecture*

The system follows a **multi-agent orchestrator pattern** where specialized components handle distinct responsibilities while maintaining loose coupling. The architecture comprises four primary layers:

- **Conversation Layer:** Manages user interactions and dialogue flow
- **Recommendation Engine:** Processes queries and retrieves relevant suggestions
- **Data Layer:** Handles vector embeddings, user clustering, and data storage
- **Orchestration Layer:** Coordinates workflows across all system components

## Technology Stack

### *Large Language Model Integration:*

- OpenAI GPT-4o-mini for intent classification, slot extraction, and contextual reranking
- Temperature-controlled responses (0.1-0.4) for consistency
- Rate limiting implementation for API call management

### *Vector Database Infrastructure:*

- ChromaDB for distributed vector storage and similarity search
- Sharded architecture supporting horizontal scaling across 8 shards
- HuggingFace sentence-transformers/all-MiniLM-L6-v2 for text embeddings
- Support for 1M+ document capacity with 150k documents per shard

### *Machine Learning Pipeline:*

- K-means++ clustering algorithm for user segmentation
- Scikit-learn pipeline with 10-dimensional feature space
- User profiling based on demographics, preferences, and behavioral patterns
- Silhouette score optimization for cluster validation

## **Detailed Component Analysis**

### *Conversation Layer:*

#### **OpenAIConversationAgent** (conversation\_agent.py):

The primary orchestrator managing dialogue flow through systematic state progression. Implements continuous conversation flow without terminating sessions, maintaining context across multiple turns.

Key features:

- Dual slot update mechanisms: new\_query for complete context replacement, slot\_updating for incremental updates
- Integration with query enhancement and vector retrieval pipelines
- Comprehensive error handling with fallback responses
- Question history tracking and conversation state management

#### **OpenAIIntentClassifier** (intent\_classifier.py):

Utilizes OpenAI GPT-4o-mini to classify user intents across seven categories:

- RECOMMEND, FILTER\_UPDATE, CLARIFICATION, FEEDBACK, GREETING, GOODBYE, OTHER
- Fallback pattern matching for API failures
- Confidence scoring with threshold-based decision making

**ConversationMemory** (memory.py):

Stateful memory management system tracking dialogue progression:

- Slot-based information storage with 7 predefined slots (dietary, cuisine\_1, cuisine\_2, item\_name, price, meal\_type, label)
- Context preservation mechanisms for maintaining conversation coherence
- Turn-by-turn history with timestamp tracking and slot update logging

**ResponseGenerator** (response\_generator.py):

Systematic response generation following a clear progression pattern:

- State-driven dialogue flow: item preferences → dietary requirements → budget → search confirmation
- JSON-structured response parsing with fallback mechanisms
- Question deduplication to prevent repetitive interactions

**Slot Extraction System** (slot\_extract.py):

Advanced entity extraction combining OpenAI LLM capabilities with rule-based fallbacks:

- Intent-aware slot filling with context preservation
- Support for cuisine mapping across 70+ cuisine types
- Flavor enhancement and dietary restriction handling
- Null value management preventing generic responses

## *Recommendation Engine*

**OpenAIQueryEnhancer** (query\_enhancer.py):

LLM-powered query refinement system that transforms user preferences into optimized search queries:

- Semantic query construction combining cuisine types, specific dishes, and preferences
- ChromaDB-compatible filter generation with complex boolean logic
- Price handling with strict, approximate, and range conditions
- Cuisine mapping across comprehensive taxonomy of food types

**ShardedRetrievalAgent** (shards\_retrieval.py):

Distributed vector search system managing queries across multiple ChromaDB shards:

- Parallel shard querying with configurable top-k results per shard
- Automated result aggregation and deduplication
- Support for complex metadata filtering

## **Two-Stage Contextual Reranker** (rerank.py):

Advanced reranking system using dual-stage LLM evaluation:

### *Stage 1: Context Analysis*

- Temporal context analysis examining eating progression patterns
- Document inventory mapping with cuisine categorization
- Adaptive ranking condition generation based on user journey
- Comprehensive document evaluation with mandatory reasoning

### *Stage 2: Final Ranking*

- Top-10 document selection based on contextual scoring
- Multi-criteria evaluation (critical, high, medium, low priority)
- Quality assurance validation for ranking coherence
- Fallback ranking using rating-price ratio for API failures

## *Data Layer*

### **Embeddings Setup** (embeddings.py):

Vector embedding infrastructure optimized for CPU deployment:

- sentence-transformers/all-MiniLM-L6-v2 model with normalized embeddings
- Memory-aware configuration based on available system resources
- CPU optimization for cost-effective cloud deployment

### **User Clustering Agent** (user\_clustering\_agent.py):

Machine learning model for user segmentation and personalization:

- K-means++ clustering with 4 user personas
- 10-dimensional feature space: Age, Monthly\_Income, Gender, Weather, Marital\_Status, Average\_Rating, Average\_Z\_Score, purchase\_sensitivity, C\_Type, Veg\_Ratio
- Silhouette score of 0.279 indicating reasonable cluster separation
- Default value handling for missing user attributes

## *Orchestration Layer*

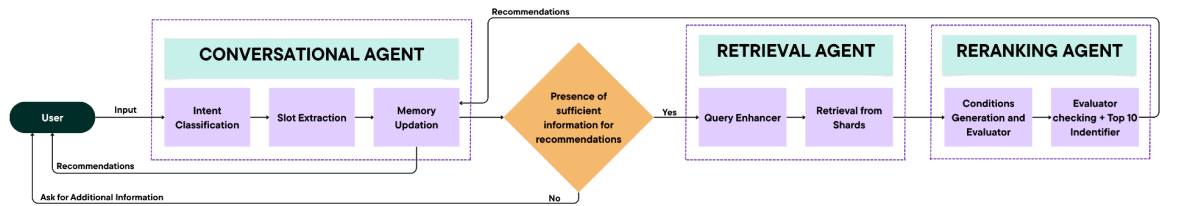
### **RecommenderOrchestrator** (orchestrator.py)

Central coordination hub managing end-to-end recommendation workflow:

- Multi-agent coordination with shared memory state
- Configuration-driven component initialization
- Recommendation context handling with history tracking

# Technical Implementation Details

## Data Flow Architecture



The system processes user requests through a sophisticated pipeline:

1. **Input Processing:** User utterances are captured and normalized
2. **Intent Classification:** OpenAI determines user intent with confidence scoring
3. **Slot Extraction:** Entities are extracted and stored in conversation memory
4. **Query Enhancement:** LLM refines search queries with semantic optimization
5. **Vector Retrieval:** Distributed search across sharded ChromaDB instances
6. **Contextual Reranking:** Two-stage evaluation for relevance optimization
7. **Response Generation:** Systematic dialogue progression with context awareness

## Vector Database Sharding Strategy

The system implements horizontal partitioning across 8 ChromaDB shards:

- **Shard Distribution:** 150,000 documents per shard.
- **Collection Strategy:** Individual collections per shard with persistent storage
- **Query Parallelization:** Concurrent querying across all shards with result aggregation
- **Fault Tolerance:** Individual shard failure handling without system-wide impact

## Memory Management and State Persistence

### Conversation State Tracking:

- Slot-based information storage with null value initialization
- Context-aware updates preserving existing information
- Turn-by-turn history with comprehensive metadata
- Session persistence across conversation boundaries

### User Profile Integration:

- Demographic and behavioral feature extraction
- Real-time cluster assignment for new users

- Preference learning from conversation context
- Historical interaction tracking

## *LLM Integration Patterns*

### **Multi-Temperature Strategy:**

- Intent classification: 0.1 for consistency
- Slot extraction: 0.2 for balanced creativity
- Response generation: 0.4 for natural language variety
- Reranking: 0.1 for deterministic evaluation

### **Rate Limiting and Error Handling:**

- 60 requests per minute throttling
- Exponential backoff for API failures
- Fallback mechanisms for each LLM-dependent component
- Graceful degradation maintaining system availability

## *Reranking Algorithm Deep Dive*

The two-stage reranking system implements sophisticated contextual evaluation:

### **Stage 1 Analysis Framework:**

- **Temporal Context:** Analyzes time gaps and eating progression patterns
- **Document Inventory:** Maps cuisine types, price ranges, and quality tiers
- **Condition Generation:** Creates measurable ranking criteria based on user journey
- **Comprehensive Evaluation:** Provides reasoning for every retrieved document

### **Stage 2 Ranking Logic:**

- **Critical Conditions:** Exclusionary criteria for top-3 rankings
- **High Priority:** Creates clear ranking tiers
- **Tie-Breaking:** Uses medium/low conditions for fine-tuning
- **Quality Assurance:** Validates ranking logic and journey alignment

# Data Processing and Feature Engineering

## *User Clustering Model*

The system employs a sophisticated user segmentation approach:

### **Feature Engineering:**

- Age and income standardization
- Behavioral metrics (rating patterns, purchase sensitivity)
- Demographic encoding (gender, marital status, weather preference)
- Cuisine type preferences and vegetarian ratios

### **Clustering Pipeline:**

- K-means++ initialization for stable clustering
- 4-cluster solution with balanced distribution
- Silhouette analysis for cluster validation
- Persona mapping for interpretability

### **User Personas Identified:**

1. Established Urban Professionals
2. Premium Self-Employed Segment
3. Young Urban Students
4. Price-Sensitive Employees

## *Vector Embedding Strategy*

### **Document Processing:**

- Text normalization and preprocessing
- Metadata extraction and enrichment
- Semantic embedding generation using sentence transformers
- Vector storage with comprehensive metadata

### **Query Processing:**

- Real-time embedding generation for user queries
- Semantic similarity computation using cosine distance
- Multi-field search across content and metadata
- Result scoring and relevance ranking

# System Integration and Workflow

## *End-to-End Recommendation Flow*

### 1. **User Profile Initialization:**

- Demographic data collection
- Cluster assignment using trained ML model
- Preference initialization in conversation memory

### 2. **Conversation Management:**

- Natural language intent detection
- Progressive slot filling with context preservation
- Systematic dialogue state progression
- Multi-turn conversation handling

### 3. **Query Processing:**

- Semantic query construction from conversation context
- LLM-enhanced query refinement
- Complex filter generation for vector search

### 4. **Retrieval and Ranking:**

- Distributed vector search across sharded database
- Contextual reranking with temporal analysis
- Top-10 recommendation selection

### 5. **Response Delivery:**

- Recommendation presentation with reasoning
- Follow-up question generation
- Conversation continuity management

## Performance Characteristics

### *Retrieval Performance*

- **Vector Search:** Sub-second response times across 1M+ documents
- **Sharded Architecture:** Linear scalability with shard addition
- **Caching Strategy:** Embedding caching for frequently accessed queries



## *LLM Integration Efficiency*

- **Request Optimization:** Batch processing where possible
- **Context Management:** Efficient prompt construction and token usage
- **Fallback Mechanisms:** Reduced API dependency through local fallbacks

## *Memory and Resource Usage*

- **CPU-Optimized:** Designed for cost-effective deployment
- **Memory Management:** Efficient embedding storage and retrieval
- **Scalable Architecture:** Horizontal scaling capabilities

## **Quality Assurance and Validation**

### *Recommendation Quality*

- **Two-Stage Validation:** Contextual analysis and ranking verification
- **User Journey Alignment:** Recommendations match conversation flow
- **Diversity Metrics:** Balanced recommendation sets across categories

### *Conversation Quality*

- **Intent Classification Accuracy:** Multi-tier validation with fallbacks
- **Slot Filling Precision:** Entity extraction with confidence scoring
- **Response Relevance:** Context-aware response generation

### *System Reliability*

- **Error Handling:** Comprehensive exception management
- **Graceful Degradation:** Partial functionality during component failures
- **Recovery Mechanisms:** Automatic retry and fallback systems

This technical architecture represents a comprehensive approach to building an AI-powered food recommendation system that balances sophistication with practical deployment considerations. The modular design, advanced ML techniques, and robust error handling create a production-ready system capable of delivering personalized food recommendations through natural conversation.